

Bounded-Authority Agricultural Advisory: Selective Resolution and Evidence-Bound Verification for Safe Bengali Decision Support

Raiyaan Reza^{a,*}, Md. Golam Rabiul Alam^{b,**}

^aDepartment of Computer Science and Engineering, North South University, Dhaka 1229, Bangladesh

^bDepartment of Computer Science and Engineering, BRAC University, Dhaka 1212, Bangladesh

ARTICLE INFO

Keywords:

Agricultural decision support
Large language models
Retrieval-augmented generation
Safety-critical verification
Bounded factual authority
Selective resolution
Low-resource NLP
Edge computing

ABSTRACT

Large Language Models and standard Retrieval-Augmented Generation (RAG) pipelines are fluent in agricultural extension but carry no intrinsic bound on factual authority. In crop protection, a fluent sentence can silently misbind a pesticide to the wrong crop, inflate a dosage, or drop a mandatory Pre-Harvest Interval (PHI) — failures that remain confident-sounding and are therefore dangerous. We present the **Bounded-Authority Advisory Architecture (BAA)**, which separates factual authority from linguistic realization: the generative model may phrase and explain, but it may not invent or mutate safety-critical agronomic values. BAA enforces an 11-slot, single-record, fail-closed certification contract and selectively routes queries among deterministic resolution, grounded generation, clarification, abstention, and escalation.

On 10,000 adversarial relational-misbinding attacks, BAA rejects 100.0% of corrupted chemical claims (95% Wilson CI [99.96, 100.0]), against 80.0% dangerous acceptance for a lexical baseline and 38.65% for an LLM-as-a-judge verifier. On a 100-case live end-to-end benchmark, BAA attains 97.0% certified advisory correctness and 0.0% critical unsafe acceptance (CI [0.0, 3.7]), confirmed by a three-agronomist validation study (4.82/5 mean correctness; 100.0% chemical-safety pass; Gwet's AC1 = 0.862). Calibrated abstention sustains 84.56% advisory coverage at a 1.26% selective-risk level across 20,112 held-out cases. The deterministic resolution ladder eliminates LLM invocation for 61.5% of queries, cutting weighted latency 2.28× and per-query cost by 89.7%. The results demonstrate that subordinating generative language to bounded factual authority yields a dependable, auditable, and affordable foundation for safety-critical agricultural decision support.

1. Introduction

1.1. Agricultural Advisory as a Safety-Critical Computing Problem

A wrong agricultural recommendation can remain fluent, grounded-looking, and dangerous. In crop protection, a recommendation is not a paragraph of prose; it is a tuple of decision-critical values—which crop, which pest or pathogen, which registered active ingredient, which formulation, which dosage interval, which unit and solvent volume, which spray interval, and which pre-harvest interval (PHI). Consider a recommendation for potato late blight: the correct answer specifies mancozeb 80 WP at 2.0–2.5 g/L water, sprayed every 7 days, with a 14-day PHI. Changing any single field—the crop, the dose, the interval—creates a different recommendation with different consequences. An error in any of these fields can cause phytotoxicity, crop loss, chemical residues on food, environmental contamination, or poisoning of the applicator.

Large Language Models (LLMs) are fluent in all of these fields. Retrieval-Augmented Generation (RAG) systems [12, 14] can make them superficially grounded as well. Yet neither fluency nor retrieved text determines the correctness of a safety-critical agricultural value. The decisive

question is not *can the model produce a plausible sentence?* but *who holds factual authority over the recommendation?* In unconstrained and RAG-based systems, that authority implicitly rests with the generative model's parameters and with whatever text happens to be retrieved.

1.2. Why Retrieval Alone Is Insufficient

Retrieval quality in agricultural domains varies sharply with query type and language condition. Text-first retrieval Hit@1 on standard formal Bengali is roughly 72–73%, but drops to roughly 45–48% on regional dialects and 41–42% on romanized Banglish input (Layer E17, threshold sweep at $\theta=0.7$). The failure is not solved by better embeddings alone: when retrieved documents individually contain valid facts from different records, a generator will assemble a semantically plausible composite—the correct crop paired with a dose taken from a different pathogen's record, or the correct active ingredient paired with an uncalibrated PHI from an unrelated row.

These failures are the dominant failure mode, not rare edge cases. In a 100-case root-cause audit (E10), retrieval evidence omission accounted for 28% of cases and generative relational misbinding—the generator pairing a host crop with a chemical from an adjacent row—for 22%. Both are silent: the output remains fluent and confident, which is precisely why they are dangerous.

*Corresponding author

**Co-corresponding supervisor

raiyaan.reza@northsouth.edu (R. Reza);

rabiul.alam@northsouth.edu (Md.G.R. Alam)

ORCID(s): 0009-0002-8610-8673 (R. Reza); 0000-0002-8610-8673 (Md.G.R. Alam)

1.3. The Missing Architectural Question

Existing agricultural advisory systems fall into two camps. Generative RAG systems (Farmer.Chat [14], KrishokBondhu [1]) treat the LLM as the end-to-end reasoning engine. Structured knowledge systems (decision tables, knowledge graphs, crop simulation wrappers) hold facts but cannot converse with farmers in natural language [8, 9]. A growing body of work addresses individual reliability properties—temporal validity [11], multimodal diagnosis [5], post-generation faithfulness [6], and RAG security [10]—but no evaluated system makes the authority relationship between deterministic agricultural facts and generative language explicit and evaluates the safety–coverage–latency frontier of that relationship.

1.4. Proposed Solution: Bounded-Authority Selective Resolution

We propose the **Bounded-Authority Advisory Architecture (BAA)**, a decision-support framework for safety-critical crop protection. Its central principle is an explicit design rule:

Design Rule. In safety-critical agricultural advisory, the component that generates fluent language must not be the component that holds factual authority.

BAA assigns factual authority to structured, versioned, provenance-carrying evidence records. The LLM is retained as a *linguistic realization* component: it may phrase, explain, and disambiguate, but it may not invent or mutate the values of the safety-critical contract. Throughout this paper, **BAA** refers to the architecture; **KrishokChat** refers to the deployed system artifact. Each advisory is certified by an 11-slot single-record binding predicate, and when certification cannot be established, the system selectively routes to deterministic resolution, grounded generation with verification, conversational clarification, safe abstention, or escalation to the national Krishi Call Center (16123).

1.5. Research Questions

1.6. Contributions

1. **Bounded-authority advisory architecture.** A decision-support architecture that separates factual authority from generative language through an 11-slot, fail-closed, single-record certification contract (Section 4).
2. **Risk-/evidence-aware selective resolution.** A five-tier resolution ladder that routes each query among deterministic resolution, grounded generation, clarification, abstention, and escalation, formalized as a decision policy over evidence and risk state (Section 3).
3. **Multi-axis empirical evaluation.** An evidence chain covering adversarial safety, agronomist-validated correctness, risk–coverage calibration, linguistic and temporal robustness, injection security, and constrained-deployment economics, across 36 experimental layers and $\approx 60,000$ evaluation cases (Section 6 and Table 5).

Table 1

Research questions, key metrics, and result sections.

RQ	Question	Key Metric	Section
RQ1	Does BAA certify correct, safe advice on naturalistic and adversarial queries, as judged by expert agronomists?	CAC, Gwet's AC1	§7
RQ2	Does single-record joint binding protect against misbinding, counterfactual corruption, and evidence conflict?	CUAR, Rejection Rate	§8
RQ3	Can calibrated abstention trade coverage for a user-specified risk bound?	AURC, Coverage@Risk	§9
RQ4	Does the architecture preserve safety under linguistic register, temporal, and multimodal uncertainty?	Dialect Coverage, Gazette Adherence	§9
RQ5	What are the latency, cost, connectivity, and channel-fidelity properties of the constrained pathways?	Latency, Zero-LLM%, Cost/1k	§10

The remainder of this paper is organized as follows. Section 2 positions the work against the literature. Section 3 formalizes the problem and the certification contract. Section 4 describes the architecture, and Section 5 the knowledge governance. Section 6 presents the methodology, Sections 7–10 the results organized by research question, and Sections 11–15 the discussion, limitations, reproducibility, and conclusion.

2. Related Work

2.1. Agricultural Conversational AI and RAG Advisory

The deployment of LLM-based agricultural extension has grown rapidly. Farmer.Chat [14] demonstrated scalable multilingual conversational advisory for smallholder farmers, engaging over 15,000 farmers across four countries. KrishokBondhu [1] showed a Bengali voice-enabled retrieval pipeline. These systems establish that conversational interfaces are usable in low-resource agricultural settings, but they treat the LLM as the end-to-end reasoning engine. Under dialectal variation and multi-document retrieval, they remain exposed to the silent factual corruption documented in Section 1. Our positioning is complementary: we do not claim a new conversation interface; we evaluate whether an explicit authority layer can make such interfaces safe for safety-critical crop-protection advice.

2.2. Structured Agricultural Knowledge and Domain-First RAG

A parallel line of work embeds structured agricultural knowledge into decision support. The goat-farming knowledge assistant of Han et al. [8] textualizes decision trees and tables to improve answer consistency over heterogeneous livestock knowledge. TARAG [11] contributes a time-aware knowledge base and time-sensitive re-ranking so that crop-pest recommendations respect phenology and pest life

stages. SMART [5] integrates structured multimodal embeddings with human-in-the-loop verification for plant disease diagnosis. DSSAT-LM [9] couples LLMs with mechanistic crop simulation by parsing free text into structured simulation parameters. These works establish that structured representations and time awareness improve agricultural advisory quality. BAA incorporates the same mechanisms—structured typed facts, temporal validity tags, and multimodal conditioning—but makes them *enforced constraints* of a certification contract rather than retrieval features.

2.3. RAG Faithfulness, Verification, and Security

Faithfulness of generated answers to retrieved evidence has become a first-class evaluation target. RAGChecker [12] argues that retrieval and generation must be diagnosed separately. Generate-but-Verify [6] formalizes post-generation faithfulness prediction as an explicit task with precision/recall variants and shows that better faithfulness assessment improves downstream RAG performance. FaithfulRAG [15] models fact-level conflict between retrieved knowledge and parametric memory, and claim-level provenance tracing is demonstrated in scientific verification by SciTrue [2]. On the security side, SafeRAG [10] demonstrates that manipulated or poisoned retrieved knowledge can defeat otherwise strong RAG pipelines. These works motivate our architecture: they show that post-hoc textual verification and retrieval-side fixes are necessary but insufficient for safety-critical fields, because they verify *language* rather than *relational authority*. BAA differs by certifying an 11-field tuple against a single evidence record, so a faithful-sounding sentence whose fields come from different records fails certification by construction.

2.4. Selective Answering, Abstention, and Risk–Coverage

Selective prediction and abstention are established tools in classification. Risk–coverage curves and calibrated abstention allow a system to trade coverage for a bounded error rate. In conversational systems, abstention is increasingly treated as a first-class output rather than a failure [6]. BAA applies these tools to agricultural advisory: abstention is certified *appropriate* when the 11-slot contract cannot be jointly satisfied, and the abstention threshold is frozen on a development split and validated on a held-out test split (Section 9).

2.5. Gap and Positioning

Table 2 summarizes the positioning. Existing systems exhibit strong individual properties—multilingual fluency, structured facts, temporal awareness, multimodality, offline operation—but no evaluated system binds a safety-critical advisory contract to a single authoritative record and then selectively resolves across deterministic, generative, and abstention paths. The novelty we claim is not any individual component; it is the *intersection*: a reliability-oriented agricultural computing architecture in which factual authority is explicitly bounded, uncertainty can terminate safely, and generative language is subordinate to validated evidence.

3. Problem Formulation and Design Requirements

3.1. Advisory State and Decision Space

Let an incoming farmer interaction be denoted $\mathbf{x} = \langle q_t, I_v, \mathcal{M}, \mathcal{H} \rangle$, where q_t is the natural-language query (standard Bengali, a regional dialect, or romanized Banglish), I_v is an optional uploaded crop photograph, $\mathcal{M}$ is environmental metadata, and $\mathcal{H}$ is conversation history.

The system maps $\mathbf{x}$ to an action $a(\mathbf{x}) \in \mathcal{A}$:

$$\mathcal{A} = \left\{ \begin{array}{ll} \text{CERTIFY,} & \text{deterministic resolution,} \\ \text{GEN_VERIFY,} & \text{grounded verification,} \\ \text{CLARIFY,} & \text{disambiguation,} \\ \text{ABSTAIN,} & \text{fail-closed refusal,} \\ \text{ESCALATE,} & \text{human handoff (16123)} \end{array} \right\} \quad (1)$$

ESCALATE routes deterministically to the Bangladesh government’s Krishi Call Center (dial 16123), a nationally active agricultural helpline staffed by certified agronomists.

3.2. The Critical Agricultural Tuple (11-Slot Contract)

In crop protection, an actionable recommendation is an 11-field tuple (the agronomic decision contract) C :

$$C = \langle c, p, s, a, f, d_{\min}, d_{\max}, u, v_s, t_{\text{int}}, t_{\text{phi}} \rangle \quad (2)$$

where c is host crop, p pest/pathogen, s growth stage, a active ingredient, f formulation, $[d_{\min}, d_{\max}]$ dosage interval, u unit, v_s water volume denominator, t_{int} spray interval, and t_{phi} pre-harvest interval (PHI).

3.3. Evidence Authority Model

Let $\mathcal{E}$ be the set of typed evidence records in the fact store. Each record $e \in \mathcal{E}$ carries two distinct attributes:

$$\text{Authority}(e) = g(\text{source, jurisdiction})$$

$$\text{Validity}(e) = g(\text{date, version, regulatory state})$$

A record is *authoritative* if it originates from an accredited institution (MoA, BARI, BRRI, DAE) for Bangladesh. A record is *current* if its effective date is unexpired and regulatory status is active. Both must hold simultaneously.

The certification predicate (Figure 1) is the central constraint of the architecture:

$$\text{Certify}(C) = 1 \iff \exists e^* \in \mathcal{E} : \left[\begin{array}{l} \text{Authoritative}(e^*) \wedge \\ \text{Current}(e^*) \wedge \\ \text{JointlyEntails}(e^*, C) \wedge \\ \text{Complete}(C) \wedge \\ g(\mathbf{x}) \geq \theta^* \end{array} \right] \quad (3)$$

Reading each conjunct in plain terms: (i) $\text{Authoritative}(e^*)$ requires the evidence to originate from a nationally accredited source; (ii) $\text{Current}(e^*)$ requires the record’s effective/expiry dates and regulatory polarity to be valid at

Table 2

Systematic positioning of KrishokChat alongside contemporary agricultural AI, retrieval-augmented decision support, and low-resource advisory systems.

System / Framework	Language	Domain	RAG	Structured Facts	Temporal	Multimodal	Fail-Closed	Edge/Offline
Farmer.Chat [14]	Multilingual	General Agri	✓	Partial	×	×	×	×
KrishokBondhu [1]	Bengali	Crop Advisory	✓	×	×	Voice	×	×
Domain-First RAG [8]	German/English	Livestock	✓	✓	×	×	×	×
TARAG [11]	Chinese/English	Crop Diseases	✓	×	✓	×	×	×
SMART [5]	English/Hindi	Plant Pathology	✓	✓	×	✓	HITL	×
AgroTutor [13]	Spanish	Crop Planning	×	✓	×	×	×	✓
KrishokChat (This Work)	Bengali	Crop Advisory	✓	✓	✓	✓	✓	✓

query time; (iii) $\text{JointlyEntails}(e^*, C)$ requires all 11 slots to be satisfied by the *same* record—records may never be assembled across sources; (iv) $\text{Complete}(C)$ requires every safety-critical slot to be non-null; and (v) $g(\mathbf{x}) \geq \theta^*$ is the calibrated selective score exceeding the frozen abstention threshold. If any conjunct fails, the system fails closed: it must not complete the recommendation; it clarifies, abstains, or escalates. This implements the design principle that a failed safety-critical evidence condition is never repaired by generating more text.

3.4. Primary Endpoints

We use the following primary endpoints throughout, reported with 95% Wilson confidence intervals where sample sizes permit. All endpoints were pre-registered before test-set evaluation.

1. **CUAR** — Critical Unsafe Acceptance Rate: *unsafe cases certified as safe / total unsafe evaluated cases*. Primary safety endpoint.
2. **CAC** — Certified Advisory Correctness: *correct certified advisories / certified advisories issued*. Primary advisory endpoint.
3. **Resolution / advisory / safe certified coverage**: respectively the fraction of queries receiving some output, the fraction of answerable queries receiving useful advice, and the fraction of answerable queries receiving a *correct certified* advisory.
4. **AA** — Appropriate Abstention: *correct abstentions / cases requiring abstention*.
5. Per-slot accuracy for each of the 11 contract fields, reported individually.

3.5. Reliability Requirements

Table 3 maps each reliability requirement to its failure mode and the fail-closed response required of the architecture. This table is the design contract used in Sections 4–5 and the evaluation in Section 6.

4. Bounded-Authority Advisory Architecture (BAA)

4.1. Architectural Overview

BAA is organized as a five-tier selective resolution ladder with a downstream relational verifier (Figure 2). The

design principle is that resolution and verification are deterministic and typed, while generation is a bounded realization step.

Tier 0 — Risk and context gate. Lightweight deterministic screening for banned/restricted chemicals, poisoning and self-harm framings, and prompt-injection payloads. Matches carry a regulatory polarity check before any retrieval or generation occurs.

Tier 1 — Capability and perception router. Entity normalization across Bengali registers (standard, dialectal, Banglish) plus optional crop classification from an uploaded image. Detection metadata narrows the retrieval and resolution search space.

Tier 2 — Structured fact authority and deterministic resolver. Typed, versioned, provenance-carrying records in a SQLite fact store. A resolver maps (crop, problem, stage, intent) to a fact record and renders a certified answer from a validated template. No LLM is invoked on this path.

Tier 3 — Grounded generative realization. For queries with evidence but no deterministic path, an LLM phrases an answer under a single-record constraint contract: it may explain the retrieved facts in Bengali but may not alter numerical or regulatory fields.

Tier 4 — Verification and fail-closed output. The relational verifier re-derives the candidate tuple from the (possibly generated) response and checks it against the single evidence record. Any slot discrepancy revokes certification, suppressing the response and routing to CLARIFY, ABSTAIN, or ESCALATE.

4.2. Risk/Context Gate

The gate applies a deterministic, six-way policy classification before any retrieval: *safe_agri*, *banned_or_restricted_chemical*, *self_harm_or_poisoning_risk*, *off_topic*, *prompt_injection*, and *low_confidence*. Terminal classes (banned chemical, self-harm/poisoning, injection) return canned safe responses that redirect to the national Krishi Call Center (dial 16123) without invoking the LLM. This implements the fail-closed property at the earliest point: a safety-class query can never reach generation.

4.3. Capability Router and Entity Normalization

A lightweight character n -gram intent classifier (1.25 MB, p50 latency 0.3855 ms; Layer E25) extracts crop, pest, and intent-type metadata at negligible cost, with a joint

Table 3

System reliability requirements, operational failure modes, and bounded-authority containment mechanisms.

Reliability Requirement	Potential Failure Mode	Advisory Risk	System Response / Containment
Evidence Completeness	Missing PHI / Dosage bounds	Chemical residue / Toxic crop harvest	Fail-closed abstention (ABSTAIN)
Relational Integrity	Cross-record pesticide misbinding	Ineffective or phytotoxic application	Single-record joint binding certification
Visual Semantic Alignment	Crop/Disease classification error	Misdirected agronomic advice	Metadata routing fallback & clarification
Perception Uncertainty	Low confidence visual input ($\tau < 0.80$)	Erroneous disease identification	Fallback to conversational clarification
Temporal Validity	Obsolete chemical recommendation	Application of withdrawn/banned agrochemicals	Temporal invalidation & policy override
Adversarial Robustness	Prompt injection / Delimiter hijack	System instruction override	Deterministic structured extraction filter
Channel Bandwidth Limit	SMS truncation / Buffer overflow	Silent loss of vital safety warnings	Deterministic 160-char GSM template
Network Degradation	High packet loss / Cloud disconnect	Service unavailability in rural edge	Offline SQLite cryptographic cache

exact match of 78.4% (95% CI [72.89, 83.05]). Because routing metadata is imperfect, BAA does not rely on it for authority: the verifier and the certification predicate (Eq. 3) absorb routing uncertainty, so a wrong metadata label cannot produce an unsafe certification. When image metadata is available, the candidate search space collapses from 2,135 nodes to a mean of 516.4 nodes (-75.64% ; Layer E17, threshold 0.8), with a misrouting rate of 3.15% (95% CI [2.65, 3.74]) that is contained rather than amplified by the downstream certification.

4.4. Structured Fact Authority and Deterministic Resolution

The fact store holds typed records with a 20-field schema (crop, problem, stage, active ingredient, formulation, dose bounds, unit, volume, interval, PHI, regulatory status, source ID, source version, effective/expiry dates, verifier identity, evidence hash). The deterministic resolver traverses the fact graph: at hop depth 3, traversal yields 100% of the 11 contract slots with 100% provenance traceability in 0.0128 ms mean latency (Layer E19, 23-node proof-of-mechanism graph). We emphasize that this is a proof of mechanism on a small graph, not a scalability claim.

4.5. Generative Realization with Constrained Decoding

When the deterministic path cannot fire but evidence exists, Tier 3 invokes the LLM (a fine-tuned Gemma-4 4-bit model in the local deployment, or an API model in the cloud variant) with the constraint: you may explain these facts in Bengali; you must not alter numerical or regulatory fields. Generation is always followed by the verifier. The architecture deliberately does not ask the LLM to infer missing evidence; missing evidence routes to recovery (retrieval, normalization, structured retry) or abstention.

4.6. Relational Certification and Fail-Closed Execution

The verifier checks: schema validity, numerical consistency against the record's dosage bounds, unit consistency, formulation compatibility, crop–problem compatibility, temporal validity, source binding (the 11-slot tuple must come from one record hash), and banned/restricted registry status. Certification is issued only when all checks pass against a single record. This is the mechanism that makes cross-record and counterfactual misbinding mathematically

impossible, and it is the mechanism empirically evaluated in Section 8.

4.7. Temporal and Provenance Controls

Each record carries effective/expiry dates and an authority manifest (Section 5). At certification time, the resolver validates the record against the manifest; a superseded record fails the Current(e) predicate and cannot be certified even if its text is retrieved. Every certified output emits a provenance chain (answer claim $\rightarrow$ structured fact $\rightarrow$ source node $\rightarrow$ source document $\rightarrow$ document version/date).

4.8. Offline, Edge, and Channel Support

A signed offline fact pack (SQLite) supports deterministic resolution and verification with zero network access. The cache is updated via differential hash-chained deltas (Layer E23) so that a disconnected client can detect tampering and receive incremental updates. This has a direct practical implication for deployment in Bangladesh: a farmer on a *char* (river island) with no mobile connectivity can still receive a certified pesticide advisory, resolved entirely on-device from the offline fact pack. Output channels (web, mobile, SMS) render from the certified structured record, so constrained channels re-encode the contract rather than re-generate safety-critical content.

5. Knowledge Base Construction and Governance

5.1. Institutional Corpus and Source Provenance

The factual foundation of KrishokChat is a curated corpus of 2,946 institutional source documents published by accredited Bangladeshi agricultural institutions, covering 47 field crops, 312 pest/disease entries across 9 agro-ecological zones, and 189 active ingredient profiles with regulatory polarity tags:

- **BARI Krishi Projukti Hatboi (9th Edition, 2023):** crop technology manuals covering field crops, tubers, and horticultural varieties [3].
- **BRRI Adhunik Dhaner Chas (22nd Edition, 2024):** canonical rice production and pathology guidelines [4].
- **DAE/MoA agrochemical governance registries:** current lists of approved, restricted, and banned pesticides in Bangladesh (regulatory polarity source).

Every node in the fact store carries provenance: source document, source version, effective date, expiry date, verifier identity, and a cryptographic evidence hash.

5.2. The 11-Slot Evidence Schema and Authority Stratification

As detailed in Table 4, the knowledge base decouples the agronomic decision contract (Layer A: the 11 semantic slots) from the evidence governance layer (Layer B: source authority, publication/effective dates, regulatory polarity, provenance hash, verifier signature). Sources are stratified into a strict authority hierarchy:

MoA Regulatory > Institutional (BARI/BIRRI)
> Legacy Institutional > Secondary (4)

For example, if a 2024 MoA gazette bans an active ingredient that a 2021 BARI manual still recommends, the MoA record takes precedence: the older BARI record fails the Authority(*e*) and Current(*e*) predicates simultaneously and cannot be certified. When multiple sources overlap, the resolver enforces precedence by authority rank then by publication timestamp. A record must be both authoritative (by source/jurisdiction) and current (by date/version/regulatory state) to satisfy Eq. (3).

5.3. Cryptographic Hash-Chaining and Tamper-Evident Packs

Structured facts are compiled into an offline SQLite fact pack for edge and offline deployment. Each record is hashed with SHA-256 and incorporated into a Merkle-style chain (each record hash is incorporated as a leaf of an incremental hash tree; the root hash changes detectably on any leaf mutation). Updates are distributed as differential delta packs rather than full downloads. In Layer E23, 1,000 single-bit and swap mutations were injected into the chain; the client detected 100.0% of them (95% CI [99.62, 100.0]), with a delta payload of 791 bytes versus a 10,985-byte full pack (92.8% bandwidth reduction) and sub-millisecond chain verification (p95 0.1919 ms). We frame this as a deployment-integrity result, not a cryptography contribution: it ensures the offline evidence base cannot be silently altered or replayed.

5.4. Knowledge Ingestion, Maintenance, and Scalability

No model output is ever written directly into the fact store. The ingestion pipeline is a controlled chain: authoritative source → document staging → fact extraction → schema validation → conflict detection → expert review → versioning/hash/signature → release to the fact store.

Extracting the 11 agronomic slots from unstructured government PDFs requires semi-automated regex/pattern

extractors coupled with agronomist double-blind validation, taking approximately 18–25 person-minutes per crop-disease profile. The automated schema compiler rejects invalid or out-of-bound numerical ranges, preventing typographical errors from entering the hash-chain. Discrete packaging formats (e.g., “1 sachet (5 g) per 10 L water”) are handled natively via the solvent volume denominator (v_s) and formulation (f) fields, without requiring continuous float modeling.

When recurring queries fail at Tier 4 because a knowledge gap exists, the failures are aggregated into a prioritized authoring queue. Layer E24 evaluated the marginal cost of closing one gap: a targeted authoring session for chili anthracnose (BARI Krishi Projukti Hatboi 2021, p. 148) added two validated facts in 18 minutes, closing the cluster gap completely (55/55 queries, 100.0%), lifting system coverage from 58.7% to 64.2% (95% CI [61.18, 67.11])—a 5.5 pp absolute gain at 3.06 newly resolvable queries per authoring minute. We report this as a single-cluster case study, not a productivity distribution. Crucially, expanding the system to new crops or neighboring agro-climatic zones (e.g., West Bengal, Assam) requires only compiling the new jurisdiction’s official pesticide gazette into the 11-slot relational schema; no neural model retraining is required.

6. Experimental Methodology

6.1. Experimental Principles

All evaluations follow five fixed controls. **(C1) Frozen partitions:** each test set is frozen on disk before any tuning; thresholds and prompt strategies are selected on development splits only and never on test. **(C2) Identical inputs across systems:** every baseline receives the same raw query, image, and context. **(C3) Deterministic seeds:** seeded runs use seed 20260827 (or 20260813 in the first-tier layers); the exact seed per layer is recorded in the experiment registry. **(C4) Pre-registered metrics:** endpoints are the CUAR, CAC, coverage, and appropriate-abstention definitions of Section 3, plus per-layer auxiliary metrics. **(C5) Statistical protocol:** binary proportions use 95% Wilson intervals; paired comparisons use McNemar’s test with Holm–Bonferroni correction across the baseline ladder; latency is reported as p50/p95/p99 when a distribution exists; risk-coverage uses AURC and fixed-risk coverage.

6.2. Benchmark Suites

The battery spans nine evaluation corpora, comprising approximately 60,000 evaluation cases across 36 completed layers:

1. **Adversarial misbinding suite** (10,000 cases): 10 corruption families × 1,000 cases each, juxtaposing valid facts across records (Layer E02).
2. **11-slot hazard suite** (13 configurations × 10,000 cases): single-slot ablation isolating safety per contract field (Layer E03).
3. **Counterfactual binding suite** (2,000 pairs): true vs. mutated evidence certification (Layer E05).

Table 4

Two-layer evidence schema: the agronomic decision contract (Layer A, 11 semantic slots) and the evidence governance layer (Layer B).

Schema Layer	Field Name	Data Type / Format	Agronomic / Governance Purpose
Layer A: Decision Contract	crop	String (canonical Bengali)	Host crop identity (e.g., potato)
	problem	String (canonical Bengali)	Target pest/pathogen identity (e.g., late blight)
	stage	Categorical	Permitted phenological growth stage
	active	ISO-standard string	Authorized active ingredient (e.g., mancozeb)
	formulation	Standard code	Commercial formulation (e.g., 80 WP, 250 EC)
	dose_min	Float	Minimum effective dose
	dose_max	Float	Maximum safe dose
	unit	SI unit string	Dose unit (g, ml, kg)
	volume	Float + unit	Solvent volume denominator (per 10 L)
	interval	Integer (days)	Minimum spray interval
	phi	Integer (days)	Pre-harvest interval
Layer B: Evidence Governance	source_authority	Categorical	MoA / BARI / BRR / secondary (rank)
	publication_date	ISO-8601	Release date of the institutional manual
	effective_date	ISO-8601	Administrative effective date
	regulatory_polarity	Enum (APPROVED / BANNED)	National regulatory status
	provenance_hash	SHA-256 hex	Digest of source text and node metadata
	verifier_signature	Ed25519	Institutional release signature

- Multi-register benchmark** (4,000 cases): Standard Bengali, farmer queries, dialects, and Banglish, 1,000 each (Layers E06, E17, E21).
- Risk-coverage corpus** (20,112 cases): 12,068 train / 4,022 dev / 4,022 test; threshold frozen on dev (Layer E04).
- Adversarial security suite** (1,400 cases): seven attack families including native and Banglish injections (Layers E07/E08, E22).
- Temporal conflict suite** (100 cases/system): obsolete-vs-current gazette and source conflict tests (Layer E30).
- Cross-modal conflict suite** (100 cases/system): text crop identity vs. visual classification (Layer E31).
- Rural network suite** (1,000 queries $\times$ 4 profiles): 4G/3G/Edge/2G latency, loss, and jitter profiles (Layer E14).

6.3. Research Question Mapping and Layer Nomenclature

To help navigate the experimental battery, Table 5 maps each experiment tag to its governing Research Question and primary evaluation target. Readers may use this table as a roadmap before diving into individual result sections.

6.4. Baseline Ladder and Fairness Controls

Because the underlying failure modes differ per layer, we define a canonical baseline ladder and instantiate the relevant subset in each layer. All systems share the same frozen inputs within a layer (C2). Baselines were configured using standard published industry best-practices—standard LangChain/LlamaIndex BM25+dense RAG with official hyperparameters; standard zero-shot and few-shot evaluation prompts for LLM-as-a-judge. The LLM judge baseline was evaluated under standard off-the-shelf single-turn verification prompts and was not specifically prompt-engineered with multi-step temporal scratchpads, reflecting

standard production guardrail deployments rather than an optimized upper bound.

- **B0 — Unconstrained LLM.** Direct generation (GPT-4o-Mini in live-API layers).
- **B1 — Vanilla RAG.** BM25/lexical retrieval feeding the generator without authority constraints (Llama-3.1-8B in live layers).
- **B2 — Guarded RAG.** Input risk-classifier before retrieval.
- **B3 — RAG + post-hoc LLM judge.** A secondary LLM evaluates textual faithfulness of the generated answer.
- **B4 — Deterministic structured resolver.** Fact traversal with no generation.
- **B5 — Evidence-constrained RAG.** Multi-document citation with a partial slot matcher but *without* single-record joint binding (the isolation baseline for the contribution).
- **B6 — BAA (this work).** The full bounded-authority architecture.

Table 6 records the mapping per layer and the exact models used.

6.5. Expert Evaluation Protocol

For the agronomist study (Layer E13), three certified agricultural extension specialists rated 200 advisory outputs (50 per system) on correctness (1–5 Likert), chemical safety pass, evidence traceability, and deployment approval. Outputs were presented without system labels in a blinded, randomized order; raters assigned system IDs only post-hoc. Agreement is reported with Gwet’s AC1 first-order agreement coefficient [7]; the primary coefficient on the safety

Table 5

Taxonomy of Experimental Layers Mapped to Research Questions (RQ1–RQ5).

Research Question	Core Investigation Target	Evaluation Layers	Key Empirical Metric(s)
RQ1: Advisory Quality	End-to-end correctness, expert review, dialects	E06, E13, E17, E21, E27	CAC (%), Gwet's AC1, Dialect Recall
RQ2: Safety & Authority	Adversarial misbinding, metaphoric, conflicts	E02, E03, E05, E07, E08, E11, E28–E31, E38, E39	CUAR (%), Mutation Rejection (%), CBC
RQ3: Selective Reliability	Risk–coverage calibration, failure taxonomy	E04, E10, E25, E34, E37	AURC, ECE, Coverage at Fixed Risk
RQ4: Robustness	Bengali registers, temporal validity, cross-modal	E12, E17, E21, E22, E23, E24, E30, E35	Gazette Adherence (%), Register Coverage
RQ5: Systems Efficiency	Low-connectivity, latency, zero-LLM economics	E09, E14, E15, E18, E19, E20	Latency (ms), Zero-LLM (%), Cost/1k (\$)

Table 6

Baseline ladder used across layers. Not every layer runs every system; each layer table states the exact systems and models used. B6 is the contribution.

ID	Architecture	Operational Workflow	Factual Authority
B0	Unconstrained LLM	Direct generation, no retrieval	Parametric memory
B1	Vanilla RAG	Lexical/BM25 retrieval → generation	Retrieved text
B2	Guarded RAG	Input risk classifier → RAG → generation	Retrieved text + gate
B3	RAG + post-hoc LLM judge	Generation → secondary-LLM faithfulness check	Secondary LLM
B4	Deterministic structured resolver	Fact-store traversal, no generation	Structured fact store
B5	Evidence-constrained RAG	Multi-document citation, partial slot matching, <i>no</i> single-record binding	Multi-record lexical
B6	Bounded-Authority Advisory (BAA)	Selective resolution ladder + 11-slot single-record certification	Single validated record

dimension is 0.862, with 0.814 on deployment approval and 0.785 on the correctness scale.

6.6. Reproducibility Protocol

Every layer records: script path, spec path, frozen seed, environment (OS, Python, packages), raw-output path, and where applicable the git commit of the code at run time. Layers E02–E26 carry a full verification block (self-checks, determinism re-run on 10% of cases, trace check, real-application check); the newer live-API layers (E27–E39) do not yet carry that block, and we state this explicitly in Section 13 rather than implying a uniform audit trail.

7. Results: End-to-End Advisory Quality and Expert Validation (RQ1)

BAA achieves **97.0% Certified Advisory Correctness** on live farmer queries, confirmed by three independent agronomists with a 100.0% chemical-safety pass rate. This section establishes the headline: the architecture works on naturalistic and adversarial queries, across registers, as judged by domain experts.

7.1. Live End-to-End Benchmark (E27)

BAA achieves **97.0% CAC** (95% Wilson CI [91.55, 98.97]) versus 74.0% (B0 GPT-4o-Mini), 73.0% (B1 Llama-3.1-8B), and 58.0% (B4 LLM-judge guardrail). On the primary safety endpoint, BAA records **0.0% critical unsafe acceptance** (CUAR, 95% CI [0.0, 3.7]) across all 100 cases, whereas B0 accepted unsafe recommendations in 8.0% of cases, B1 in 6.0%, and the LLM-judge guardrail (B4) in

Table 7End-to-end advisory performance on the live benchmark (Layer E27, $n = 100$ per system; 70 naturalistic + 30 adversarial Bengali queries; 100% real API calls).

Baseline	CAC (%)	CUAR (%)	Abstain (%)	Latency (ms)
B0: GPT-4o-Mini	74.0 [64.6, 81.6]	8.0 [4.1, 15.0]	14.0 [8.5, 22.1]	2,726.3
B1: BM25 (Llama-3.1)	73.0 [63.6, 80.7]	6.0 [2.8, 12.5]	5.0 [2.2, 11.2]	6,010.8
B4: LLM Judge	58.0 [48.2, 67.2]	15.0 [9.3, 23.3]	16.0 [10.1, 24.4]	1,597.0
B6: BAA (Ours)	97.0 [91.6, 99.0]	0.0 [0.0, 3.7]	33.0 [24.6, 42.7]	3.8 (mod.)

15.0%—worse than the unguarded baseline. Table 7 reports the full comparison.

Two framing caveats are explicit. First, BAA abstains at 33.0% (CI [24.56, 42.69])—more than double any baseline—so the 97.0% CAC is computed over the answered/certified set, not raw accuracy. Second, the BAA latency entry (3.8 ms) is a modeled in-process figure, not a measured wall-clock percentile; the measured decomposition is reported in Section 10.¹

7.2. Agronomist Validation (E13)

Three certified agronomists rated 200 advisory outputs. BAA scored a mean correctness of **4.82/5** (SD 0.28), **100.0% chemical-safety pass rate** (95% CI [98.12, 100.0]), 98.5% evidence traceability, and 96.5% deployment approval. RAG+LLM judge scored 4.10/5 and 86.0% safety pass; vanilla RAG 3.65/5 and 71.5%; unconstrained LLM 3.12/5 and 64.5%. LLM-direct and vanilla RAG received

¹†3.8 ms is the modeled verification-and-render latency from the E9 stage decomposition; all live baselines report measured p50/p95 from real API calls.

deployment approval in only 38.0% and 52.5% of cases. The raters' agreement on safety was Gwet's AC1 = 0.862, with 0.814 on deployment approval and 0.785 on correctness scale.

7.3. Multi-Register Correctness (E06)

Across 4,000 multi-register queries (1,000 per register), BAA certified 65.8% of cases correctly and abstained safely in 34.2%, with **0.0% dangerous acceptance** (95% CI [0.0, 0.1]) in all four registers. Per register: standard formal Bengali 76.2% correct; authentic farmer 68.3%; regional dialects 59.5%; romanized Banglish 59.2%. The strongest baselines in the same suite—vanilla RAG—reached only 56.07% overall correct with 13.45% dangerous acceptance (CI [12.43, 14.54]); direct LLM 58.73% with 1.85%. The register invariance of the BAA safety boundary is the robustness headline of Section 9.

7.4. Per-Contract-Slot Correctness

The 11-slot contract is the unit of correctness. Removing any single slot from the certification contract raises dangerous acceptance from 0.0% to between 1.9 pp (spray interval) and 31.6 pp (dosage bounds); removing all typed constraints collapses the system to the 80.0% hazard rate of a lexical substring matcher. The per-slot contribution is reported in full in Section 8 (Layer E03 ablation).

End-to-end correctness establishes that BAA works on naturalistic queries, expert-rated. The harder question is whether it can be fooled. Section 8 answers that.

8. Results: Safety Verification and Authority Enforcement (RQ2)

BAA rejected **100.0% of 10,000 adversarial relational-misbinding attacks** (95% CI [99.96, 100.0]) and 100.0% of 11,000 metamorphic mutations. The safety boundary holds across every attack family evaluated. This section traces the structural mechanism behind that result and stress-tests it from every angle.

Two exploratory layers (E32, oracle-vs-predicted routing; E33, high-confidence wrong metadata) recorded non-zero CUAR values for BAA (11.0% and 10.0%). These layers tested routing metadata perturbation, a scenario where the CUAR definition does not apply identically to evidence certification; both are reported as exploratory and are discussed openly in Section 13.

8.1. Relational Misbinding Defence (E02)

Layer E02 evaluated 10,000 adversarial cases (10 corruption families $\times$ 1,000 cases), each presenting juxtaposed valid facts from different records designed to trigger relational misbinding. BAA rejected **100.0%** (0/10,000 false certifications; 95% CI [99.96, 100.0]). A lexical substring matcher certified 80.0% of corrupted claims as safe (CI [79.20, 80.77]); dense retrieval 37.56%; citation-alignment 43.14%; LLM-as-judge 38.65%; vanilla RAG 72.65%; the

partial 8-slot matcher (B5) 24.5%. Single-record joint binding makes cross-record misbinding impossible by construction: a claim whose 11 slots cannot all be satisfied from a single accredited record hash fails certification regardless of how the slots individually appear in the retrieved pool. Table 8 reports the full comparison.

8.2. Metamorphic Authority Testing (E28)

Layer E28 generalizes the attack: 11 mutation operators (crop, pest, active ingredient, formulation, dosage overdose, dosage unit, water volume, interval shortening, PHI shortening, regulatory polarity flip, provenance-hash corruption) $\times$ 1,000 cases each = 11,000 adversarial mutations, run against seven verification systems. BAA rejected **100.0%** (95% CI [99.97, 100.0]) with 0.0% false certification. The partial 8-slot matcher (B5) rejected 63.64% overall but **0.0%** on exactly the four safety-critical families it lacked (water volume, application interval, PHI, provenance hash)—demonstrating that partial slot coverage is qualitatively different from full coverage, not merely a degraded version. The LLM-as-judge rejected only 72.25%, the citation-alignment baseline 57.33%, dense matcher 37.05%, BM25 36.36%, and unconstrained LLM 18.41% (false-certifying 81.59%).

8.3. Counterfactual Evidence Binding (E05)

Layer E05 tested certification under counterfactual mutations (dose scaled 5 $\times$, PHI shortened to 3 days): 2,000 evidence pairs per system. BAA certified 0.0% of mutated claims (95% CI [0.0, 0.19]) while certifying 100.0% of true evidence, giving a Counterfactual Binding Consistency (CBC) of 1.0000. Vanilla RAG certified 72.65% of mutated claims (CI [70.65, 74.56]; CBC 0.2735), a lexical matcher 59.55% (CBC 0.4045), and an LLM-as-judge 38.65% (CBC 0.5705). The verifier is evidence-bound rather than prior-bound.

8.4. Parametric-vs-Evidence Conflict (E29)

In Layer E29, each of five LLMs was run in four modes: direct generation, standard RAG, prompted judge, and BAA. When retrieved evidence conflicted with the model's parametric prior, direct generation showed parametric intrusion rates of 30–56%, standard RAG still leaked 8–30%, and a prompted judge leaked 2–16%. The BAA mode achieved **100.0% evidence adherence and 0.0% parametric intrusion** in every model/mode cell, because the deterministic 11-slot resolver makes parametric memory irrelevant: it certifies the record when the contract holds and abstains otherwise. Note the effective sample granularity caveat reported in Section 13.

8.5. Temporal Source Authority (E30)

BAA achieved **100.0% current-gazette adherence** (95% CI [96.30, 100.0]) with 0.0% obsolete leakage (CI [0.0, 3.7]) in the temporal conflict benchmark (100 cases per system). The unconstrained LLM cited the obsolete record in 39.0% of cases; lexical mixed RAG in 39.0%; and the LLM-judge guardrail in **75.0%**. The mechanism is structural: the Current(e) predicate is evaluated against the authority

Table 8

Comparative authority evaluation under metamorphic mutation (Layer E28, 11 operators $\times$ 1,000 cases = 11,000 mutations per system). BAA (B6) is the only system reaching 100.0% rejection; the partial 8-slot matcher (B5) drops to 0.0% on exactly the four safety-critical slot families it lacks (volume, interval, PHI, provenance hash). B0/B4 latencies are modeled values, not measured wall-clock.

Verification Strategy	Rejection Rate (%)	False Certification (%)	Observed CUAR (%)	Latency p95 (ms)
B0: Unconstrained LLM	18.41 [17.70, 19.14]	81.59 [80.86, 82.30]	67.15 [66.26, 68.02]	1,450.0
B1: Lexical BM25 Matcher	36.36 [35.47, 37.27]	63.64 [62.73, 64.53]	45.45 [44.53, 46.39]	0.85
B2: Dense Embedding Matcher	37.05 [36.15, 37.95]	62.95 [62.05, 63.85]	46.47 [45.54, 47.41]	8.40
B3: Citation Multi-Doc Alignment	57.33 [56.40, 58.25]	42.67 [41.75, 43.60]	32.96 [32.09, 33.85]	12.40
B4: LLM-as-a-Judge Guardrail	72.25 [71.41, 73.08]	27.75 [26.92, 28.59]	21.84 [21.07, 22.62]	1,680.0
B5: Partial 8-Slot Matcher	63.64 [62.73, 64.53]	36.36 [35.47, 37.27]	27.27 [26.45, 28.11]	1.20
B6: 11-Slot Single-Record BAA (Ours)	100.0 [99.97, 100.0]	0.0 [0.0, 0.03]	0.0 [0.0, 0.03]	3.8

manifest at resolution time, so no model temperature or knowledge-cutoff randomness can override a deterministic authority comparison.

8.6. Cross-Modal Perception Uncertainty (E31)

When text and visual crop identity conflicted, BAA triggered a safe clarification request in **100.0%** of conflicted inputs (95% CI [98.77, 100.0]) and delivered 0.0% wrong chemical recommendations. The unconstrained multimodal LLM delivered the wrong recommendation in 54.0% of cases. On the harm axis, BAA's advantage over the LLM-judge is not in unsafe output (both 0.0%) but in the clarification rate (100.0 vs 26.0%). Per the project's artifact constraints, perception inputs are classification labels, not detection bounding boxes.

8.7. Slot Ablation: What Each Contract Field Buys (E03)

Layer E03 ablated each of the 11 contract slots individually (10,000 cases per configuration, 13 configurations). Removing each slot raises dangerous acceptance as follows: dosage bounds +31.6 pp, regulatory polarity +17.8 pp, active ingredient +14.2 pp, host crop +11.4 pp, target pathogen +8.2 pp, solvent volume +7.1 pp, formulation +5.8 pp, unit +5.4 pp, growth stage +4.9 pp, PHI +2.7 pp, spray interval +1.9 pp. Removing all typed constraints collapses the system to the 80.0% hazard of a lexical substring matcher. The dosage-bound and regulatory-polarity slots are individually load-bearing: their ablation alone creates the two largest hazard surges in the battery.

The safety boundary holds under attack. But a system that refuses everything is safe yet useless. Section 9 asks: can BAA calibrate when to abstain?

9. Results: Selective Resolution, Calibration, and Robustness (RQ3, RQ4)

Calibrated abstention sustains **84.56% advisory coverage at a 1.26% selective-risk level** across 20,112 held-out cases—a better coverage-risk trade-off than all four alternatives. Across linguistic registers, temporal conflicts, and retrieval degradation, the safety boundary (0.0% CUAR, 95% CI [0.0, 3.7]) holds in every evaluated cell.

Table 9

Risk-coverage calibration on held-out test split (Layer E04, 20,112 cases; 12,068 train / 4,022 dev / 4,022 test).

Selective Policy	AURC ($\downarrow$)	ECE ($\downarrow$)	Coverage (%)	Risk (%)
Raw Generator Conf.	0.0761	0.0896	10.22	1.70
Lexical Overlap Score	0.1124	0.0855	0.37	6.67
LLM-Judge Conf.	0.0310	0.1066	36.60	0.54
Conformal Baseline	0.0179	0.1231	77.62	1.06
Calibrated BAA (Ours)	0.0153	0.0785	84.56	1.26

9.1. Risk-Coverage Calibration (E04)

Layer E04 evaluated five selective policies on a 20,112-case corpus (12,068 train / 4,022 dev / 4,022 test). The threshold θ^* was frozen on the development split and transferred to the held-out test split. Table 9 reports test-split results.

The calibrated relational policy achieves the lowest test Area Under the Risk-Coverage curve (AURC, **0.0153**) and Expected Calibration Error (ECE, **0.0785**) of all five policies. At the frozen threshold $\theta^* = 0.2375$, it sustains **84.56% test coverage at a selective risk of 1.26%**—higher coverage at comparable risk than the conformal baseline (77.62% at 1.06%). At fixed risk budgets, raw generator confidence achieves 0% coverage and the relational policy 84.56% at comparable risk levels.

9.2. The Lightweight Intent Router (E25)

The character n -gram router achieves a joint exact match of 78.4% (95% CI [72.89, 83.05]) with per-field crop accuracy 96.4% and intent-type accuracy 100.0%, at p50 latency 0.3855 ms and a memory footprint of 1.25 MB. We present this as a routing-efficiency result, not a routing-accuracy claim: the 78.4% joint-EM figure motivates the design—the architecture absorbs routing uncertainty in the certification predicate rather than trusting metadata.

9.3. Root-Cause Failure Taxonomy (E10)

A manual audit of 100 residual failure/abstention cases assigned each to one of eight categories (Table 10): retrieval evidence omission 28%, generative relational misbinding 22%, numerical/volumetric parsing ambiguity 14%, dialectal lexical ambiguity 12%, regulatory status discrepancy

Table 10

Taxonomy of residual failure modes from 100-case root-cause audit (Layer E10), architectural detectors, and mitigation policies.

Failure Root Cause	Share (%)	Observed Mechanism	Architectural Detector	Mitigation Strategy
Retrieval Omission	28%	Lexical/dense search misses relevant node	Coverage threshold θ_{cov}	Fallback to MD-chunk search (R13)
Relational Misbinding	22%	Multi-document context conflation	11-slot relational verifier	Single-record joint binding certification
Numerical/Volume Ambiguity	14%	Missing water volume or concentration unit	Unit/volume slot extractor	Automated structured clarification
Linguistic Ambiguity	12%	Colloquial slang or under-specified symptom	Register normalizer / Intent gate	Multi-turn disambiguation prompt
Regulatory Discrepancy	8%	Conflicting state vs trade advice	Regulatory authority ranker	Enforcement of MoA banned-pesticide list
Temporal Obsolescence	6%	Superseded BARI manual cited	Metadata publication date filter	Cache invalidation & version pinning
Verifier Conservatism	6%	Valid advice rejected due to strict bounds	Threshold calibration (θ^*)	Continuous human-in-the-loop review
Calibration Underconfidence	4%	High-quality generation scored below θ	Conformal risk regularizer	Risk-coverage curve optimization

8%, temporal/source expiry 6%, conservative verifier over-refusal 6%, and calibration underconfidence 4%. Critically, both dominant failure families were converted into fail-closed abstention rather than dangerous advice: the verifier caught all 22 misbinding attempts (0.0% dangerous leak) and the coverage threshold converted the 28 omissions into abstentions.

9.4. Architectural Ablation: Where the Safety Comes From (E34)

Layer E34 isolates each architectural layer with a cumulative ablation (100 cases per tier). The base LLM (L0) certifies 61.0% correctly with 14.0% CUAR. Adding lexical retrieval (L1) raises correctness to 72.0% and cuts CUAR to 4.0%; hybrid dense retrieval (L2) reaches 72.0% at 98.0% operational coverage but still 3.0% CUAR. The prompted safety guardrail (L3) worsens three of four metrics (62.0% correctness, 10.0% CUAR, 78.0% coverage), consistent with E27 and E30 findings. Adding the SQLite fact base with the five-tier ladder (L4) drives correctness to 100.0% and CUAR to 0.0% at 70.0% certified coverage; the full BAA (L5) matches L4 on every accuracy metric. We state the honest reading: L4-to-L5 does not itself demonstrate the verifier's safety contribution—that evidence lives in the E28 B5-vs-B6 contrast—and the coverage drop from 98% to 70% is larger than the correctness gain from 72% to 100%, so the ablation is a coverage-for-certainty trade, not a monotonic improvement.

9.5. Input Robustness: Bengali Register Coverage (E17, E21)

Bengali input spans at least four registers: standard formal, authentic farmer language, regional dialects (Sylheti/Chittagonian), and romanized Banglish. Text-first retrieval degrades sharply on the last three: Hit@1 falls from 72.1–73.4% (standard) to 44.7–45.7% (dialects) and 41.5–41.6% (Banglish) at thresholds 0.7–0.8 (E17). Detection-gated routing collapses the candidate search space by 75.64% (2,135 $\rightarrow$ 516.4 nodes) and lifts Hit@1 on dialects by +36.6 pp and on Banglish by +39.7 pp. Layer E21 then asks what the routing choice does to end-to-end advisory outcomes per register: coverage rises from 42.9% to 89.8% on dialects (+46.9 pp), and from 41.6% to 90.0% on Banglish (+48.4 pp), while holding hazard at **0.0% in all eight register $\times$ routing cells** (0/4,000; 95% CI [0.0, 0.09]). Table 11 summarizes the register-level outcomes.

Table 11Register-robust routing outcomes (Layer E21, $n = 1,000$ per register $\times$ 2 routing paths; 95% Wilson CIs).

Register	Routing Path	Coverage (%)	Abstain (%)	Hazard (%)
Standard Bengali	Text-first	70.3 [67.4, 73.1]	29.7	0.0 [0.0, 0.4]
Standard Bengali	Detection-gated	93.0 [91.3, 94.4]	7.0	0.0 [0.0, 0.4]
Authentic Farmer	Text-first	58.2 [55.1, 61.2]	41.8	0.0 [0.0, 0.4]
Authentic Farmer	Detection-gated	92.6 [90.8, 94.1]	7.4	0.0 [0.0, 0.4]
Regional Dialects	Text-first	42.9 [39.9, 46.0]	57.1	0.0 [0.0, 0.4]
Regional Dialects	Detection-gated	89.8 [87.8, 91.5]	10.2	0.0 [0.0, 0.4]
Romanized Banglish	Text-first	41.6 [38.6, 44.7]	58.4	0.0 [0.0, 0.4]
Romanized Banglish	Detection-gated	90.0 [88.0, 91.7]	10.0	0.0 [0.0, 0.4]

9.6. Retrieval-Degradation Safe Degradation (E12)

Under four retrieval regimes (gold, noisy, poisoned-contradictory, omitted evidence; 2,000 queries each), vanilla RAG's chemical hazard escalates from 10.0% (gold) to 33.2% (poisoned). BAA holds **0.0% hazard in all four regimes** (each CI [0.0, 0.19]) and degrades monotonically in coverage: 84.56% $\rightarrow$ 54.3% $\rightarrow$ 0% $\rightarrow$ 0%, converting ungrounded context into abstention. This is the cleanest demonstration of safe degradation under evidence-quality uncertainty.

9.7. Temporal Validity and Knowledge Integrity (E30, E23, E24)

BAA achieved **100.0% current-gazette adherence** (95% CI [96.30, 100.0]) with 0.0% obsolete leakage in the temporal conflict benchmark (100 cases per system). The LLM-judge guardrail cited the obsolete record in **75.0%** of cases—worse than the unguarded LLM (39.0%). The mechanism is structural: the Current(e) predicate is evaluated against the authority manifest at resolution time. On tamper detection (E23), 1,000 SHA-256 chain mutations were all detected (100.0%; CI [99.62, 100.0]) at p95 0.1919 ms verification latency. Knowledge maintenance remains a data operation, not a code operation: 18 minutes closed a real cluster gap (E24, single-cluster case study).

The architecture is safe and calibrated in controlled evaluation. Can it survive rural Bangladesh—2G networks, 160-character SMS, \$50 phones? Section 10 answers that.

10. Results: Deployment Efficiency and Constrained Delivery (RQ5)

The deterministic resolution ladder eliminates LLM invocation for **61.5% of queries**, cutting weighted latency

2.28× and per-query cost by 89.7%. Deterministic SMS templates preserve 100% of safety-critical slots where LLM-composed SMS loses them in 64.4% of messages. This section reports the full systems-efficiency and constrained-channel evidence. Network profiles are simulations, not field measurements; cost figures are scenario-based on the E18 traffic mix.

10.1. LLM-Dependency Reduction Across the Resolution Ladder (E18)

Layer E18 measured the five-tier ladder's traffic mix on 5,000 modeled queries per arm. Detection-gated routing raised the zero-LLM share from 10.48% to 61.52% (95% CI [60.16, 62.86]). The decomposition: **deterministic advisory coverage** ($\tau_1+\tau_2$) is 51.04%, **safety/refusal handling** ($\tau_0+\tau_4$) is 10.48%, and **LLM-dependent traffic** (τ_3) is 38.48%. Weighted mean latency falls from 1,247.93 ms to 546.15 ms (2.28×), and serving cost from \$0.1785 to \$0.0767 per 1,000 queries. The 10.48% terminal-handling share is not “answered without LLM”; it is refused or redirected.

10.2. Stage-Level Latency Decomposition (E9)

The E9 stage decomposition (analytical, modeled) localizes the latency story: deterministic routing 0.35 ms p95, deterministic lookup 0.42 ms, hybrid retrieval 28.5 ms, verification 3.8 ms, render 0.15 ms, and generation 1,780 ms p95. Generation dominates by two orders of magnitude; verification does not. Selective routing therefore attacks the right bottleneck: every query kept out of Tier 3 avoids the generation term entirely.

10.3. Serving Cost and Economics (E20)

Per-query economics: app-online \$0.000257, SMS fallback \$0.00234, commercial cloud LLM baseline \$0.00248 (i.e., \$0.26 / \$2.34 / \$2.48 per 1,000). The app-online path is 89.65% cheaper than the commercial baseline. Assumptions: 6 queries per farmer per year, 50/30/20 app/offline/SMS mix, 120 BDT/USD; population-scale projections are deliberately excluded.

10.4. Prompt Injection and Evidence Poisoning (E07/E08)

Seven adversarial families (1,400 attacks total, 200 per family) targeted direct system override, evidence override, retrieval poisoning, delimiter hijacking, Bengali-native injection, romanized Banglish injection, and mixed code-switching. The unconstrained LLM certified unsafe output in 83.86% of attacks (95% CI [81.84, 85.69]); vanilla RAG in 69.36%; RAG with an LLM guard in 21.79% (CI [19.70, 24.02])—still vulnerable to delimiter breakouts and code-switching jailbreaks; and BAA in **0.0%** (0/1,400; CI [0.0, 0.27]), with 1,289 attacks intercepted at the deterministic Tier-0 gate and the remaining 111 blocked by the relational verifier.

Table 12

Deployment performance under constrained channels (Layers E14, E15, E18, E20, E22, E23).

Deployment Channel / Metric	Baseline	BAA (Ours)	Delta
Rural-Edge delivery @15% loss (E14)	82.0%	91.4%	+9.4 pp
Severe-2G delivery @30% loss (E14)	12.8%	58.1%	+45.3 pp
SMS critical-slot hazard (E15)	64.4%	0.0%	−64.4 pp
SMS injection survivability (E22)	36.4%	0.0%	−36.4 pp
Tamper detection rate (E23)	0.0%	100.0%	1,000/1,000
Zero-LLM workload share (E18)	10.5%	61.5%	+51.0 pp
Weighted latency (E18)	1,247.9 ms	546.2 ms	2.28×
Cost, app-online per 1k (E20)	\$2.48	\$0.26	−89.7%

10.5. Deterministic SMS Compression Fidelity (E15)

The deterministic 11-slot template preserved **100.0% of all safety-critical slots** in 102–115 characters, with 0.0% messages exceeding 160 characters and 0.0% critical hazard (95% CI [0.0, 0.38]). LLM-composed SMS fit in 92.9 characters on average but truncated or omitted safety fields: PHI survived only 35.6% of messages and the critical hazard rate was **64.4%** (CI [61.38, 67.31]). Naive truncation to 160 characters produced a 100.0% hazard rate. Once factual authority is established, constrained channels render from the certified record rather than re-generating safety-critical content.

10.6. Cellular Network Resilience (E14)

Under rural-edge loss (15%), offline-first caching sustains **91.4%** delivery versus 82.0% for cloud-only RAG (+9.4 pp); under severe 2G loss (30%), **58.1%** versus 12.8% (+45.3 pp). Cache hits resolve at p50 1.2–1.5 ms versus 8.2–11.3 s for cloud RAG. Table 12 summarizes the full deployment comparison.

10.7. Safety–Coverage–Latency Trade-off

The Pareto narrative across the paper is: BAA trades coverage (97% CAC at 33% abstention; 84.56% coverage at 1.26% risk) for safety (0.0% CUAR observed across all main evaluation layers) and wins latency on the deterministic share of traffic (2.28× weighted speedup). The trade-off is explicit and tunable through θ^* ; Figure 6 illustrates the four-panel efficiency portrait.

11. Discussion

11.1. What Bounded Factual Authority Buys

The evidence supports a single causal chain. Agricultural evidence is relational: a recommendation is only as valid as its *jointly bound* slots. Simple retrieval or citation can misbind those slots (E02: 80.0% dangerous acceptance for a lexical baseline; E05: 72.65% for vanilla RAG under counterfactual mutation). Typed single-record binding makes the misbinding impossible by construction rather than improbable by scoring (E02/E28: 100.0% rejection). Because certification is deterministic, it buys four properties beyond safety: *predictable critical fields* (the verifier checks every contract slot, and the slot ablation shows each contributes

measurably, E03), *traceability* (every certified answer carries a record hash and provenance chain), *reproducibility* (deterministic resolution is invariant by construction), and *controllable failure* (the E10 taxonomy shows residual failures convert to abstention, not to dangerous advice).

11.2. The Cost of Bounded Authority

The price is real and quantified. BAA abstains at 33.0% in the live benchmark—comparable to a human extension officer who says “I don’t know, let me check,” except BAA does it in 3.8 ms (modeled) rather than waiting for a callback. Coverage falls from 98% (hybrid RAG, at 72% correctness) to 70% certified coverage (at 100% correctness) in the E34 ablation—a larger movement than the correctness gain. The E10 audit attributes 6% of residual failures to conservative verifier over-refusal and 4% to calibration underconfidence; these are recoverable via frozen-threshold tuning but represent a standing coverage cost. Knowledge engineering is a second cost: the fact store must be maintained and versioned, though E24 shows the marginal cost is small and measurable (18 minutes to close a cluster gap).

11.3. The Role of LLMs in High-Consequence Domains

This paper is not an anti-LLM argument. The LLM remains the system’s conversational face: it handles phrasing, explanation, and disambiguation, and it is responsible for the 38.48% of traffic that the deterministic ladder cannot resolve. The empirical message is narrower and stronger: the LLM must not be the holder of factual authority for safety-critical values. The E29 parametric-conflict results show that even when a model holds the correct prior, standard RAG still leaks 8–30% of conflicting priors into output, and a prompted judge does not eliminate the leak. Structure does; prompting alone does not.

11.4. Statistical Uncertainty and Field Safety Mitigation

In the live benchmark (Layer E27, $n = 100$), BAA recorded 0.0% observed CUAR with a 95% Wilson CI of [0.0, 3.7]%. Across a national scale of 1,000,000 annual queries, an unmitigated 3.7% upper-bound rate could theoretically permit up to 37,000 hazardous recommendations. In KrishokChat, this statistical envelope is actively mitigated by three orthogonal structural defenses: (i) Tier 0 pre-retrieval screening blocks banned chemicals and restricted active ingredients in <0.1 ms before any generation; (ii) calibrated selective abstention converts ambiguous queries into clarifying questions rather than speculative completions; and (iii) structured escalation handoff routes to the national Krishi Call Center (16123) when safety-critical uncertainty cannot be deterministically resolved.

11.5. Where the Architecture Still Fails

The E10 audit is the honest failure ledger. Retrieval omission (28%) means a correct record can be missed and the query abstained on; the mitigation (dense retriever refinement, chunk-level fallback) is a coverage fix, not a

safety fix. Dialectal lexical ambiguity (12%) and numerical-unit ambiguity (14%) generate clarifications that burden the farmer. The 6% conservative-refusal rate shows the verifier can refuse agronomically plausible advice that fails a slot check—a fairness issue as much as a coverage one. Generalization beyond crop protection to veterinary dosing, human-health triage, or industrial safety represents a design-generalization argument, not an empirical one; none of those domains have been evaluated here.

12. Deployment Implications

12.1. Rural Connectivity and Offline Evidence

The E14 network results translate directly to deployment reality in Bangladesh’s rural grid: at 15% packet loss, an offline-first cache sustains 91.4% delivery versus 82.0% for cloud-only RAG, and at 30% loss the gap widens to 58.1% versus 12.8%. The offline fact pack (SHA-256-chained, 92.8% bandwidth-reduced differential updates, E23) means deterministic resolution and verification require no network at all; the network is needed only for the generative tier.

12.2. Low-End Devices

The on-device profile (E9, classification artifacts) shows INT8 ONNX models at 5.9–20.79 MB with 29–48 ms inference on mobile CPU and 100.0% parity with the FP32 reference. Together with the 1.25 MB intent router (E25) and the deterministic resolver (≤ 0.94 ms p95), the safety-critical core of the system runs entirely on commodity Android hardware; the LLM tier remains optional and cloud/API-based where available.

12.3. Extension-Service Escalation

Abstention is not abandonment: every fail-closed refusal redirects to the national Krishi Call Center (16123) or a certified agronomist, and the E38 escalation-gate evaluation shows the routing decision itself can be made deterministic (100.0% correct escalation on the corrected run, with the confounds noted in Section 8). For a national extension service, the practical implication is a division of labor: the machine handles the certifiable majority, and the human network absorbs the abstained minority at a bounded, measurable rate.

12.4. Knowledge Maintenance and Regulatory Updates

The E24 loop shows knowledge maintenance is a data operation, not a code operation: 18 minutes closed a real coverage gap. Regulatory changes (e.g., a newly banned active ingredient) are handled by the polarity registry and temporal validity fields; the E30 result shows obsolete records cannot be certified once the manifest is updated. An operator can therefore keep the system current without touching the model.

12.5. Auditability and Monitoring

Every certified answer carries a provenance chain (record hash, source, version, date) and every interaction writes an

audit event (query, risk class, route, resolution tier, evidence version, model, verification, final action, latency). This makes the system auditable per-answer — a requirement for public-sector deployment and a differentiator from chatbot-only advisory tools.

12.6. What Is Not Claimed

We have not deployed at national scale, we report simulated network profiles rather than field connectivity, and the vision modality is classification-based (task: classify) rather than detection-based. The deployment implications above are architectural and model-based, not deployment measurements.

13. Limitations and Threats to Validity

Statistical power and sample sizes. Several layers are underpowered for the effect sizes they report. The live benchmark (E27) uses $n = 100$ per system (CUAR upper CI 3.7%); the escalation layer (E38) uses $n = 30$ (upper CI 11.35%); the clarification and normalization layers (E35, E37) use $n = 45$ and $n = 37$ with control arms that saturate, so those two layers support no directional claim and are excluded from the main narrative. E29's per-cell rates are exact multiples of 2.0% despite a recorded $n = 1,000$, implying an effective granularity near $n = 50$; we report its rates without relying on its confidence intervals. All binary outcomes in the main battery use 95% Wilson intervals, but for small- n layers the intervals are wide and must be read as such. The 3.8 ms verification latency and the 546.0 ms and 1.2 ms values recur across layers with $p_{50} = p_{95}$ exactly; these are assumed constants in the results file, not measured percentiles.

Confounds in individual layers. E38 assigns a different model to each arm (GPT-4o-Mini / Llama-3.1-8B / Gemini-2.5-Flash-Lite), so the escalation headline is model-confounded; a same-model replication is required. E38 also records a post-hoc classifier fix with pre-fix numbers discarded. The E29 parametric layer's "BAA" row is invariant across modes by construction (the constraint overrides the mode), so it is one measurement, not four. E34's L4 and L5 differ only in latency/cost; the safety gain of the 11-slot verifier is established by E28's B5-vs-B6 contrast instead. Two exploratory layers (E32, oracle-vs-predicted routing; E33, high-confidence wrong metadata) recorded non-zero CUAR values for BAA (11.0% and 10.0%) that appear to contradict the 0.0% main-battery headline. As noted in Section 8, these layers tested routing metadata perturbation—a scenario where CUAR is not identically applicable to evidence certification—and both are labelled exploratory with unresolved audit status rather than hidden.

Scope: what is not measured or claimed. We have no yield, income, or adoption data. Expert agronomist ratings

(E13) measure advice quality, not field outcome. The network profiles (E14) are simulations, not field connectivity measurements; cost figures are scenario-based. The vision modality is currently whole-image classification-only; it does not perform localized bounding-box detection or spatial segmentation. Future work will integrate localized disease segmentation (SAM / YOLOv11-seg) to validate lesion coverage and severity before dosage calculation. English and other languages were not evaluated. The knowledge-graph traversal result (E19) is a proof of mechanism on a 23-node graph, not a scalability claim. The E24 maintenance result is a single-cluster case study.

Reproducibility and audit completeness. The experiment battery was produced in two phases with different audit depth. Layers E02–E26 carry a full verification block (self-checks, determinism re-run, trace check, real-application check, git commit, environment, seed). Layers E27–E39, which include the headline live-API benchmarks, do not yet carry that block in the registry; they are labelled FROZEN_AND_VERIFIED but are reproducibility-incomplete until their verification blocks are added. We state this openly rather than implying a uniform audit trail. Model and API drift, and regulatory change over time, remain ongoing validity threats that the architecture mitigates (versioned evidence, frozen thresholds) but cannot eliminate.

14. Reproducibility, Data and Code Availability

14.1. Experiment Registry

All experiments are tracked in a machine-readable registry (experiments/registry.yaml). Each layer records: research question, claim ID, spec path, runner script path, raw-output path, seed, environment (OS, Python version, package versions), timestamp, metrics block, and baseline block. Layers E02–E26 additionally record a verification block: self-checks, a determinism re-run on a 10% sample with a recorded max metric delta, a trace check, and a real-application check (backend test suite, golden replay, frontend build). Layers E27–E39 record the metrics and live-API provenance but do not yet carry the full verification block; this is stated in Section 13.

14.2. Data and Artifacts

The frozen evaluation corpora (adversarial misbinding suite, counterfactual pairs, multi-register benchmark, risk-coverage corpus, security suite, network profiles) and per-layer raw outputs are distributed with the paper artifacts. The institutional corpus (2,946 documents) originates from publicly released BARI, BRRI, and DAE/MoA publications; the derived fact store is versioned and SHA-256-pinned. The offline fact pack used in the network and cache layers is the same artifact used in production, minus any live secrets.

14.3. Configuration and Seeds

Seeds are fixed per layer (20260827 for the newer layers; 20260813 for the first-tier layers) and recorded in the

registry. The calibration threshold $\theta^* = 0.2375$ was frozen on the development split of the risk-coverage corpus and transferred unchanged to the test split; no test-set feedback was used.

14.4. Hardware and Environment

Live-API layers were run through a public model API (OpenRouter) with recorded model identifiers (GPT-4o-Mini, Llama-3.1-8B-Instruct, Gemini-2.5-Flash-Lite, Qwen-2.5-7B-Instruct, local Gemma-4 4-bit LoRA). Deterministic and local layers were executed on commodity Windows hardware (Python 3.12/3.13, pyyaml 6.0.3); the measured device profile (E9) used a mobile-class CPU with the INT8 ONNX artifacts. Model version drift is a known threat; the model identifiers above are the ones actually called at run time.

14.5. Result-Provenance Rule

Every number in this manuscript traces to a layer metrics block in the registry (the “Number Bank” extracted for this paper). No acceptance-note prose value was used where it disagreed with the tabulated metrics; the specific conflicts (E14, E17, E18, E19, E20, E21, E23, E25, E27, E28, E30, E31, E34, E35, E37, E38, E39) are documented in the Number Bank. A full audit diff is available on request.

14.6. Submission Package

The submission bundle contains: the LaTeX sources, the frozen results registry, the Number Bank extraction, all layer runner scripts and specs, the raw-output JSON files, and the figure/table generation prompts. The exact commands to regenerate any table or figure from raw output are recorded in each layer’s spec file.

15. Conclusion

We introduced and evaluated the Bounded-Authority Advisory Architecture (BAA), an agricultural decision-support architecture in which factual authority is explicitly separated from generative language realization. BAA certifies every recommendation through an 11-slot, single-record, fail-closed contract: the LLM may phrase and explain, but it may not invent or mutate safety-critical agronomic values. The architecture selectively routes each query among deterministic resolution, grounded generation with verification, clarification, abstention, and escalation to the national Krishi Call Center (16123).

The evidence shows the mechanism holds across the causal chain. Typed single-record binding rejects 100.0% of 10,000 relational misbindings and 100.0% of 11,000 metamorphic mutations, where lexical, dense, citation, and judge baselines certify 42–82% of corrupted claims. On live end-to-end evaluation, BAA certifies 97.0% correct advice with 0.0% observed critical unsafe acceptance (95% Wilson CI [0.0, 3.7]), confirmed by a three-agronomist validation study (4.82/5; Gwet’s AC1 = 0.862; 100.0% chemical-safety pass), at a 33.0% abstention cost. Calibrated abstention delivers 84.56% coverage at 1.26% risk. The deterministic resolution

ladder eliminates LLM invocation for 61.5% of queries, cutting modeled latency 2.28×; deterministic SMS preserves 100% of safety slots where LLM-composed SMS loses the PHI in 64.4% of messages; and offline-first caching sustains 91.4% delivery at 15% packet loss.

What remains unproven is equally clear. There is no farmer-outcome data, no national-scale deployment, and no same-model replication of the escalation layer. Several newer layers lack full registry verification blocks; two exploratory layers show an unresolved CUAR inconsistency; and the small- n layers cannot support the effect sizes they hint at. These define the next experimental cycle: larger independent annotation, hardened verification blocks for the live-API layers, and—ultimately—field pilots with real extension officers, where the ultimate metric is not CUAR but whether a farmer’s crop survives the season.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

The authors thank the Department of Agricultural Extension (DAE), the Bangladesh Agricultural Research Institute (BARI), and the Bangladesh Rice Research Institute (BRRI) for public access to institutional Krishi handbooks and agronomic guidance. This work is the reliability-science layer of a larger research program: prior work established the knowledge and retrieval foundations (the 1,001-query farmer benchmark and retrieval-failure analysis), and a companion systems paper describes the integrated deployment artifact. The CEA contribution is the generalizable bounded-authority architecture and its evidence chain, not the individual components, each of which exists in the literature.

Data Availability

All frozen experiment results, evaluation scripts, and the fact-store schema are available in the project repository. Layers E02–E26 carry full verification blocks; E27–E39 are reproducibility-incomplete pending verification block additions (see Section 13).

References

- [1] Ahmed, F., Kabir, H., Rahman, S., 2025. KrishokBondhu: An interactive voice-enabled bengali agricultural advisory system, in: Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP System Demonstrations), Association for Computational Linguistics. pp. 142–151.
- [2] Author, P., Author, S., 2026. SciTrue: Evidence-grounded claim verification in science, in: Proceedings of the 2026 Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations, Association for Computational Linguistics.
- [3] Bangladesh Agricultural Research Institute, 2023. Krishi projukti hatboi (handbook on agricultural technology). BARI, Gazipur, Bangladesh.

- [4] Bangladesh Rice Research Institute, 2024. Adhunik dhaner chas (modern rice cultivation). BRRI, Gazipur, Bangladesh.
- [5] Chan, Y.C., Fan, Y.C., Li, H.C., Yang, M.D., 2026. SMART: Structured multimodal agricultural retrieval-augmented transformer for plant disease diagnosis. *Computers and Electronics in Agriculture* 250, 111882. doi:10.1016/j.compag.2026.111882.
- [6] Filice, S., Haramaty, E., Horowitz, G., Karnin, Z., Lewin-Eytan, L., Shtoff, A., 2025. Generate but verify: Answering with faithfulness in RAG-based question answering, in: *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics, Asian Federation of Natural Language Processing and Association for Computational Linguistics, Mumbai, India*. pp. 1017–1037. doi:10.18653/v1/2025.ijcnlp-long.56.
- [7] Gwet, K.L., 2021. *Handbook of Inter-Rater Reliability: The Definitive Guide to Measuring the Extent of Agreement Among Raters*. 5th ed., Advanced Analytics, LLC.
- [8] Han, N., Liu, D., Norton, T., 2026. Towards an AI-based knowledge assistant for goat farmers based on retrieval-augmented generation. *Computers and Electronics in Agriculture* 245, 111542. doi:10.1016/j.compag.2026.111542.
- [9] Li, Z., Hoogenboom, G., Porter, C., Jones, J.W., 2026. Closing the gap between crop simulation and practical decision-making: DSSAT-LM semantic reasoning. *Computers and Electronics in Agriculture* 231, 109880. doi:10.1016/j.compag.2026.109880.
- [10] Liang, X., Niu, S., Li, Z., Zhang, S., Wang, H., Xiong, F., Fan, J.Z., Tang, B., Zhao, J., Yang, J., Song, S., Wang, M., 2025. SafeRAG: Benchmarking security in retrieval-augmented generation of large language model, in: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics, Vienna, Austria. pp. 4609–4631. doi:10.18653/v1/2025.acl-long.230.
- [11] Liu, L., Li, S., Qi, J., Yuan, Z., Yang, P., 2026. TARAG: A time-aware retrieval-augmented generation framework for supporting precision crop pest and disease management through large language models. *Computers and Electronics in Agriculture* 248, 111786. doi:10.1016/j.compag.2026.111786.
- [12] Ru, D., Qiu, L., Hu, X., Zhang, T., Shi, P., Chang, S., Jiayang, C., Wang, C., Sun, S., Li, H., Zhang, Z., Wang, B., Jiang, J., He, T., Wang, Z., Liu, P., Zhang, Y., Zhang, Z., 2024. RAGChecker: A fine-grained framework for diagnosing retrieval-augmented generation, in: *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, Datasets and Benchmarks Track.
- [13] Silva, J., Ortiz, R., Jarvis, A., 2020. AgroTutor: A mobile application for agricultural decision support in data-scarce environments. *Computers and Electronics in Agriculture* 175, 105580. doi:10.1016/j.compag.2020.105580.
- [14] Singh, N., Wang'ombe, J., Okanga, N., Zelenska, T., Repishti, J., Jayasankar, G.K., Mishra, S., Manokaran, R., Singh, V., Rafiq, M.I., Gandhi, R., Nambi, A., 2024. Farmer.Chat: Scaling AI-powered agricultural services for smallholder farmers. *arXiv preprint arXiv:2409.08916*.
- [15] Zhang, Q., Xiang, Z., Xiao, Y., Wang, L., Li, J., Wang, X., Su, J., 2025. FaithfulRAG: Fact-level conflict modeling for context-faithful retrieval-augmented generation, in: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics, Vienna, Austria. pp. 21863–21882. doi:10.18653/v1/2025.acl-long.1062.




Figure 1: Agricultural Claim Authority and Joint Entailment Model. An actionable recommendation must be jointly entailed by a single accredited evidence record hash. Generative LLMs participate strictly as linguistic phrasing engines without authority to alter agronomic values.




Figure 2: The Bounded-Authority Advisory Architecture (BAA). Queries enter the deterministic risk/context gate (Tier 0); the capability router normalizes language and optional image metadata (Tier 1); the structured fact authority resolves deterministically when the 11-slot contract is jointly satisfiable (Tier 2); grounded generation is permitted only under a single-record constraint (Tier 3); and the relational verifier certifies or fails closed to clarify/abstain/escalate (Tier 4).




Figure 3: Metamorphic Rejection Rates under 11,000 Systematic Corruptions per System (Layer E28). BAA achieves 100.0% rejection across all 11 mutation operators, while unconstrained LLMs accept 81.59% of corrupted advice.




Figure 4: Hazard Surge Under Contract Slot Ablation (Layer E03). Removing dosage bound enforcement yields the largest single hazard increase (+31.6 pp), while removing all typed constraints causes an 80.0% hazard collapse.




Figure 5: Risk–Coverage Curves on Held-Out Test Split (Layer E04). At the frozen threshold $\theta^* = 0.2375$, calibrated relational BAA achieves 84.56% coverage at 0.0% selective chemical risk on development criteria, with an overall test AUC of 0.0153.




Figure 6: Deployment Trade-offs: Systems Efficiency, Connectivity Resilience, and Channel Economics. (A) Delivery under rural packet loss; (B) Weighted mean latency speedup; (C) Zero-LLM resolution share (61.5%); (D) 160-character Bengali SMS slot survival rate.